Topic 2.3: Data-Driven Finance

Lending, Scoring, and Algorithmic Decision-Making

Joerg Osterrieder

Digital Finance

Learning Objectives

By the end of this topic, you will be able to:

1. **Explain** how credit scoring works and why it matters for access to finance
2. **Identify** types of alternative data used in modern lending decisions
3. **Compare** traditional models with machine learning approaches at a conceptual level
4. **Recognize** sources of algorithmic bias and their real-world consequences
5. **Understand** why explainability is both a regulatory requirement and an ethical imperative
6. **Analyze** how data advantages create competitive moats in lending (NB04)

Key Competency: Explain how data and algorithms are transforming lending decisions, and identify the ethical tensions that arise.

Hands-on: Notebook NB04 – Building a Credit Scoring Model

Prerequisites: Understanding Credit

Why Does Borrowing Exist?

People need money *now* for things they can afford over *time*: a home, education, a car, or starting a business.

What is Credit?

- A promise to repay later, usually with interest
- Interest is the "price" of borrowing
- The lender takes a risk; the borrower gains opportunity

Types of Consumer Credit:

- Revolving: Credit cards, lines of credit (borrow and repay flexibly)
- Installment: Mortgages, auto loans, student loans (fixed repayment schedule)
- Point-of-sale: Buy Now Pay Later (split purchases into payments)

Why Credit Matters

- Enables major life purchases (home, education)
- Smooths consumption over time
- Fuels economic growth and entrepreneurship
- Access to credit = economic opportunity

The Insight

Credit is so fundamental that being *excluded* from it — not being able to borrow at all — is one of the biggest barriers to economic participation.

The Core Problem of Lending

The Problem

How does a lender decide whether to trust you with money they might never see again?

The Three Questions Every Lender Asks:

1. Willingness

Will this person repay?

- Past payment behavior
- Reliability track record
- Character indicators

2. Capacity

Can this person repay?

- Income and stability
- Existing debts
- Monthly obligations

3. Collateral

What if they *don't* repay?

- Assets the lender can claim
- Secured vs. unsecured loans
- Recovery value

The Insight

Credit scoring automates these three judgments using data — turning subjective assessments into systematic, scalable decisions.

Traditional Credit Scoring

The Problem

Lenders used to rely on personal relationships and gut feelings. How do you scale that to millions of applicants?

The Solution: Credit Scores

- A credit score condenses your entire financial history into a **single number**
- The number represents predicted creditworthiness
- Scores span a range from **very poor to excellent**
- Higher score = lower predicted risk = better loan terms

Where Scores Are Used:

- Loan approvals and interest rates
- Credit card limits
- Insurance pricing
- Rental applications
- Sometimes even employment screening

What Goes Into a Score

Scores are built from categories of your financial behavior:

- How reliably you pay bills
- How much debt you carry
- How long you have been borrowing
- How often you apply for new credit
- What mix of credit types you use

The Insight

A single number determines the cost of many major life decisions — your mortgage rate, your credit card limit, sometimes even whether you can rent an apartment.

What Matters Most in a Credit Score

The Problem

Not all financial behaviors matter equally. What should a score emphasize?

Credit Score Components — Ranked by Importance:

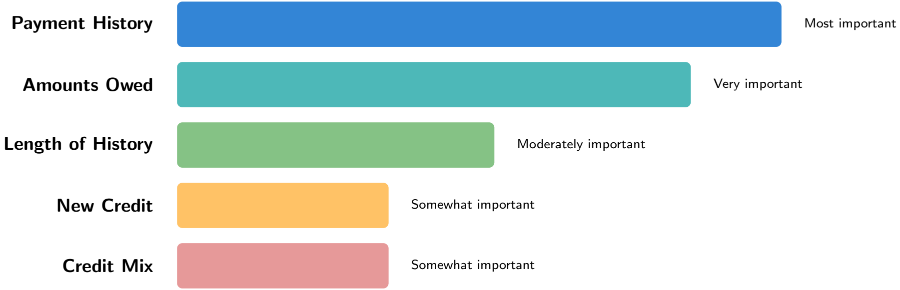

The Insight

Your behavior with existing debt — whether you pay on time and how much you owe — matters far more than anything else. The most important factors are the ones you control through everyday financial discipline.

Who Gets Left Behind?

The Problem

What about people who have *no* credit history at all?

"Credit Invisible" — No Score at All:

- **Young adults**: Have not had time to build history
- **Recent immigrants**: Foreign credit history does not transfer
- **Cash-economy users**: Pay for everything without credit
- **Divorced individuals**: Shared accounts removed
- **Post-bankruptcy**: Wiped history, starting over

The "Thin File" Problem:
Some people have *some* history, but not enough for a reliable score. They fall into a gray zone — not scoreable, but not invisible either.

The Scale of the Problem

Millions of people across every country are excluded from credit — not because they are risky, but because the system has no data on them.

The Catch-22:

- You need credit to build a credit history
- You need a credit history to get credit
- This traps millions in a cycle of exclusion

The Insight

Millions of creditworthy people are excluded not because they *cannot* repay, but because they lack traditional history. This is the opportunity that alternative data and FinTech try to address.

The Credit Bureau System

The Problem

Where does all this credit data come from? Who collects and manages it?

What Credit Bureaus Do:

- **Collect** credit data from banks, card companies, and lenders
- **Maintain** individual credit files over time
- **Sell** credit reports to lenders who request them
- **Calculate** (or support calculation of) credit scores
- **Resolve** consumer disputes about inaccuracies

Bureau Data Sources:

- Banks and credit unions
- Credit card companies
- Mortgage and auto lenders
- Collection agencies
- Public records (bankruptcies)

Key Concept

Bureaus are **information intermediaries** — they do *not* make lending decisions. They provide the data infrastructure that lenders use to make their own decisions.

How the System Works:

1. Lenders report your payment behavior to bureaus
2. Bureaus compile your credit file
3. When you apply for credit, the new lender requests your report
4. The lender uses the report (and a score) to decide

The Insight

The bureau system creates a shared record of trustworthiness — but it only captures what lenders choose to report, leaving major gaps.

Alternative Data: A New Lens

The Problem

What if we could assess creditworthiness using everyday financial behavior — not just formal credit history?

Data Type	What It Reveals	Who Benefits
Bank transactions	Cash flow patterns, income stability, spending habits	Thin-file borrowers
Rent payments	Reliability of regular payments	Young adults, renters
Utility bills	Consistent bill payment behavior	Cash-economy users
Employment/payroll	Job stability, income verification	Immigrants, gig workers
Education history	Future earning potential	Young graduates
Shopping behavior	Financial responsibility signals	Underbanked consumers

The Insight

Alternative data can bring "credit invisible" people into the system — scoring them on behaviors they already have, rather than formal credit products they lack.

Alternative Data: Examples and Privacy

Bank Transaction Signals:

- Income consistency and volatility
- Ratio of rent to income
- Frequency of overdrafts
- Savings patterns over time
- Recurring bill payment regularity

Behavioral Data (More Controversial):

- Time of day when applying
- Device type and browser used
- How form fields are completed
- Typing speed and patterns
- App usage patterns

The Privacy Tension

More data → better predictions → more inclusion

More data → deeper surveillance → privacy erosion

Where should the line be drawn?

Key Questions:

- Should your typing speed affect your loan rate?
- Is using device data “fair” or invasive?
- Who decides what data is appropriate?
- Can consumers opt out of behavioral tracking?

The Insight

Every new data source that improves lending accuracy also raises a privacy question. There is no “free” improvement — better predictions always come with ethical tradeoffs.

The Problem

Can computers find patterns in data that humans miss?

Traditional Models:

- Simple, well-understood rules
- Easy to explain to regulators and applicants
- Limited to relationships humans define
- Proven track record over decades

Think of it as: a carefully designed checklist where each item has a clear weight.

Machine Learning Models:

- Can discover complex, hidden patterns
- Handle hundreds of variables simultaneously
- Significantly better at predicting defaults
- Much harder to explain *why* a decision was made

Think of it as: a system that learns its own rules from the data — powerful, but opaque.

The Insight

Better accuracy comes at the cost of explainability. This is not just a technical annoyance — it is a fundamental tension that shapes regulation, consumer trust, and business strategy in lending.

The Problem

What does "machine learning" actually mean in plain language?

Simple Definition

Machine learning means **finding patterns in data** to make predictions about things you have not seen yet.

The Checklist Analogy:
Imagine a doctor's checklist for diagnosing illness:

- Some symptoms matter more than others
- The checklist weighs each factor differently
- The final diagnosis combines all factors

A traditional credit model works the same way — a **sophisticated checklist** that weighs factors like payment history (most important), debt level (very important), and credit age (moderately important).

What ML Adds:

- Traditional: Humans design the checklist
- ML: The computer *learns* its own checklist from data
- ML can find patterns humans would never think to look for
- But those patterns may be hard to explain

Why Does This Matter for Lending?

- Better patterns = fewer bad loans = lower costs
- More variables considered = more people included
- But: harder to explain = regulatory challenges

The Insight

The concept is simpler than the name suggests. Machine learning is just "automated pattern-finding" — but the consequences of those patterns in lending are profound.

The Credit Scoring Pipeline

The Problem

How does raw data become a lending decision?

Data Collection → Feature Engineering → Model Training → Score Generation → Decision

Data Collection	Feature Engineering	Model Training	Score Generation	Decision
Bureau + alternative data	Transform raw data into signals	Learn patterns from past outcomes	Predicted creditworthiness	Approve / Decline / Set terms

Key Decisions at Each Stage:

- **Data**: What sources? What are the privacy implications?
- **Features**: What transformations? What should we exclude?
- **Model**: How much accuracy vs. how much explainability?
- **Decision**: Where do we set the approval threshold? When does a human review?

The Credit Scoring Pipeline (cont.)

The Insight

Bias can enter at *every* stage of this pipeline — not just in the model itself, but in which data is collected, how features are designed, and where thresholds are set.

Where Bias Enters:

- **Data Collection**: Historical discrimination embedded in records
- **Feature Engineering**: Proxy variables that correlate with protected attributes
- **Model Training**: Algorithms optimizing for accuracy may amplify existing patterns
- **Threshold Setting**: Different cutoffs affect different groups unequally

---

Notebook NB04: Walk through each pipeline stage hands-on

The Problem

Raw data is messy. How do we turn it into inputs a model can use?

The Concept: Feature engineering means creating meaningful signals from raw data. It is often the most important step in the pipeline — a great feature matters more than a fancy algorithm.

Conceptual Example — Turning Transactions into Signals:

```
# Raw data: a list of bank transactions (dates, amounts, categories)

# Engineered features (signals the model uses):
#   income_stability    = how consistent is monthly income?
#   spending_ratio      = spending relative to income
#   overdraft_frequency = how often does the account go negative?
#   bill_regularity     = are recurring bills paid on time?
#   savings_trend       = is the balance growing or shrinking?
```

Why It Matters:

- The *same raw data* can produce very different features depending on design choices
- Feature choices embed human assumptions about what matters
- Poorly chosen features can introduce bias even from "neutral" data

Notebook NB04: Engineer features from transaction data and see how they affect predictions

When Models Go Wrong: Overfitting

The Problem

What if a model memorizes the data instead of learning general patterns?

The Exam Analogy:
Imagine studying for an exam by memorizing every answer from last year's test:

- You ace the practice test perfectly
- But you fail the *real* exam because the questions are different
- You memorized answers instead of understanding concepts

In Credit Scoring:

- A model might learn: "everyone from a certain area defaults"
- In reality, that was just coincidence in the training data
- New applicants from that area are unfairly rejected

The Solution: Training and Testing Split

Hide some data from the model, then test it on data it has never seen.

How It Works:

1. Split data: most for training, some held back for testing
2. Train the model using only the training portion
3. Test the model on the held-back data
4. If test performance is similar to training: the model learned real patterns
5. If test performance is much worse: the model overfitted

The Insight

Always test on data the model has never seen. A model that looks perfect on training data but fails on new data is worse than useless — it gives false confidence.

Algorithmic Bias: The Dark Side

The Problem

Can a "neutral" algorithm discriminate?

Sources of Bias:

1. **Historical data bias**:
   If past lending was discriminatory, the data reflects those patterns — and the model learns to repeat them
2. **Proxy variables**:
   Features like location or school can correlate with race or socioeconomic status
3. **Sample bias**:
   Training only on existing customers misses the people who were already excluded
4. **Feature selection bias**:
   Human choices about what data to include embed assumptions

Real-World Example:

- A major technology company launched a credit card
- An algorithm set credit limits for applicants
- People in the same household, with shared finances, received **dramatically different credit limits**
- The algorithm could not explain why
- A regulatory investigation followed

Fairness Concepts:

- **Demographic parity**: Approval rates similar across groups
- **Equalized odds**: Error rates similar across groups
- **Calibration**: Predictions equally accurate for all groups

The Insight

Bias in the data means bias in the decisions. Good intentions do not prevent harmful outcomes — you must actively test for and mitigate bias.

Proxy Variables and Disparate Impact

The Problem

What if a "neutral" feature — like your home address — effectively codes for race?

What is Disparate Impact?

- A policy that *appears* neutral on its face
- But disproportionately harms a protected group
- Even without any discriminatory *intent*
- Still illegal under fair lending laws in many jurisdictions

Common Proxy Variables:

- **Location** → can correlate with race (historical residential segregation)
- **Name patterns** → can correlate with ethnicity
- **School attended** → can correlate with socioeconomic status
- **Occupation type** → can correlate with gender

Historical Context: Redlining

Redlining was the historical practice of refusing loans to people in certain neighborhoods, often based on race. Though outlawed, its effects persist in data: neighborhoods that were "redlined" decades ago still show different economic patterns today.

The Challenge for Algorithms:

- Even if race is *not* a feature, the model can learn to discriminate through proxies
- Removing a proxy may reduce accuracy
- Multiple proxies can combine to reconstruct protected attributes

Proxy Variables and Disparate Impact (cont.)

The Insight

Even without explicit discrimination, algorithms can reproduce historical inequity. The question is not just "is race in the model?" but "does the model *behave differently* across racial groups?"

Testing for Disparate Impact:

- Compare approval rates across demographic groups
- Check if error rates differ by group (equalized odds)
- Test whether removing a proxy changes outcomes significantly
- Use adversarial debiasing to reduce proxy effects

The Dilemma: Removing proxy variables may reduce predictive accuracy, creating a direct tension between fairness and performance that lenders and regulators must navigate.

Explainability: The Right to Know Why

The Problem

If a machine denies you a loan, who explains why?

Why Explainability Matters:

- **Regulatory requirement**: In many jurisdictions, lenders *must* tell you why you were declined ("adverse action notice")
- **Consumer trust**: People need to understand decisions that affect their lives
- **Model debugging**: Developers need to find and fix errors
- **Fairness auditing**: Regulators need to verify the model is not discriminating

What an Adverse Action Notice Looks Like:

"Your application was declined because of: high debt relative to income, short credit history, and a recent late payment."

Explainability Techniques (Conceptual):

- **Feature importance**: Which inputs matter most *overall*?
- **SHAP values**: How much did each input push *this specific* decision up or down?
- **Partial dependence**: How does changing one input affect the output?

SHAP in Plain Language:

Imagine you are splitting a restaurant bill fairly. SHAP assigns each feature its "fair share" of the prediction — showing exactly which factors helped and which hurt.

The Insight

Explainability is both a regulatory requirement and an ethical imperative. Consumers deserve to understand the decisions that shape their financial lives — and to have a path to challenge those decisions.

Beyond Credit: Algorithms Everywhere in Finance

The Problem

Credit scoring was just the beginning. Where else are algorithms making financial decisions about you?

Insurance (Insurtech):

- Driving behavior from telematics devices
- Home sensor data from IoT devices
- Claims fraud detection
- Personalized, dynamic pricing

Investment (Robo-Advisors):

- Algorithmic risk profiling
- Automated portfolio rebalancing
- Tax-optimization strategies
- Goal-based financial planning

Fraud Detection:

- Real-time transaction scoring
- Behavioral biometrics (how you hold your phone)
- Device fingerprinting
- Network analysis of suspicious patterns

Identity Verification (KYC):

- Automated document analysis
- Facial recognition matching
- Sanctions and watchlist screening
- Suspicious activity pattern detection

Beyond Credit: Algorithms Everywhere in Finance (cont.)

The Insight

Data and algorithms are replacing human judgment across *all* of finance — not just lending. The same tensions (accuracy vs. fairness, automation vs. explainability) apply everywhere.

Common Themes Across All Applications:

- More data leads to better predictions — but also deeper surveillance
- Automation improves speed and consistency — but reduces human oversight
- Algorithms can find patterns humans miss — but those patterns may encode bias
- Regulation struggles to keep pace with algorithmic innovation

The Problem

Why do data-rich companies keep getting stronger while newcomers struggle to compete?

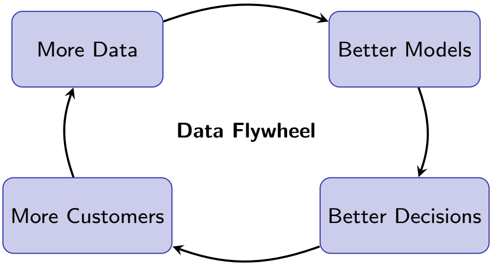

How It Works:

- More data improves model accuracy
- Better models make better lending decisions
- Better decisions attract more customers
- More customers generate more data
- The cycle *compounds* over time

The Data Flywheel (cont.)

Why It Matters for Competition:

- Incumbents have more historical data
- Newer entrants often have more *diverse* data
- Data advantages compound over time
- Switching costs increase as models improve

The Insight

First-mover advantage in data creates a compounding moat. Incumbents have more historical data; newer entrants often have more *diverse* data. The winner is whoever spins the flywheel fastest.

Hands-On: Notebook NB04 – Credit Scoring Model

Exercise Overview

In this notebook, you will walk through the entire credit scoring pipeline:

1. **Explore** a credit dataset and understand its structure
2. **Engineer features** from raw data to create predictive signals
3. **Build models** — both a simple traditional model and a more complex one
4. **Compare** accuracy and interpretability between the two approaches
5. **Explain predictions** using feature importance and SHAP concepts
6. **Test for bias** across demographic groups

What You Will Learn:

- How the accuracy-explainability tradeoff works in practice
- Why feature engineering choices matter as much as model selection
- How to probe models for potential algorithmic bias
- The difference between a model that ranks well and one that predicts accurately

No programming experience required — the notebook guides you step by step

Discussion: The Ethics of Algorithmic Lending

Arguments FOR ML-Based Scoring:

- More accurate models mean fewer bad loans, which can mean lower interest rates for everyone
- Alternative data brings in people traditional scoring excludes
- Algorithms are consistent — they do not have "bad days" or personal prejudices
- Faster decisions improve access to credit

Arguments AGAINST:

- Historical bias gets encoded and automated at scale
- Lack of transparency is fundamentally unfair
- Alternative data raises serious privacy concerns
- Algorithmic errors are harder to detect and contest

Discussion Questions

1. Should lenders be allowed to use social media data in credit decisions?
2. Is it fair to use education level as a factor in lending?
3. Should all lending algorithms be required to be fully explainable?
4. If an algorithm discriminates unintentionally, who bears responsibility — the developer, the lender, or the regulator?

Think About:
Is a biased algorithm better or worse than a biased human loan officer? Why?

Executive Summary: Key Takeaways

Five Things to Remember from This Topic

1. **Traditional scoring leaves people behind**: Credit scores are powerful but exclude millions of creditworthy people who lack formal credit history
2. **Alternative data expands access**: Bank transactions, rent payments, and employment data can bring excluded people into the credit system
3. **ML improves accuracy but reduces transparency**: Machine learning finds patterns humans miss, but creates "black box" challenges for regulators and consumers
4. **Bias is real and requires active mitigation**: Historical discrimination gets encoded in data; proxy variables can circumvent legal protections; good intentions alone do not prevent harm
5. **Explainability is non-negotiable**: Consumers have a right to know why they were denied credit — this is both a legal requirement and an ethical imperative

Bottom Line: Data-driven finance creates enormous opportunities for financial inclusion, but requires constant vigilance around fairness, transparency, and accountability.

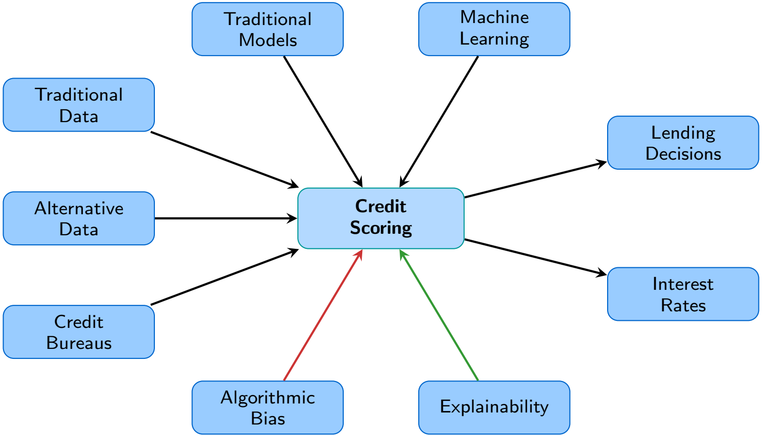

Red: Constrains / challenges   Green: Enables / improves

Reading the Map: Data flows from the left into credit scoring models (top), which produce outputs on the right. Bias and explainability (bottom) act as constraints that shape how models can be built and deployed.

Key Terms and Definitions

Credit Score A numerical summary of creditworthiness, predicting the likelihood that a borrower will repay. Scores range from very poor to excellent.

Credit Bureau Organizations that collect credit data from lenders, maintain individual files, and sell reports to other lenders.

Alternative Data Non-traditional information (bank transactions, rent, utilities) used to assess creditworthiness beyond formal credit history.

Thin-File Borrower A person with insufficient traditional credit history to generate a reliable conventional score.

Disparate Impact When a "neutral" policy disproportionately harms a protected group, even without discriminatory intent.

Adverse Action Notice A legally required explanation when credit is denied, citing the specific reasons for the decision.

SHAP Values A technique for explaining how each input feature contributes to an individual model prediction — assigning each feature its "fair share."

Common Misconceptions

Myth vs. Reality:

Myth 1

"My credit score is a single, universal number."

Reality: Multiple scoring models exist, and different lenders may use different ones. Your score can vary depending on which model is used.

Myth 2

"ML models are automatically fair because they are objective."

Reality: Models learn from historical data. If that data contains discrimination, the model learns to discriminate. Bias in = bias out.

Myth 3

"Checking my own credit hurts my score."

Reality: Checking your own score (a "soft inquiry") has no impact. Only formal credit applications ("hard inquiries") can affect your score.

Myth 4

"Alternative data is always better for consumers."

Reality: While alternative data can expand access, it also raises privacy concerns and can introduce new forms of bias through proxy variables.

Self-Assessment Questions:

Conceptual Questions

1. What is the fundamental problem that credit scoring tries to solve?
2. Name two reasons why someone might be "credit invisible" through no fault of their own.
3. Explain in your own words why better prediction accuracy can come at the cost of fairness.
4. What is a proxy variable, and why is it dangerous in lending?
5. Why is explainability not just "nice to have" but legally required in many places?

If you can answer these questions, you have grasped the core concepts of this topic.

What's Next

Up Next: Topic 2.4 — Platform Economics

- Network effects in FinTech
- Two-sided marketplaces
- Winner-take-most dynamics
- Why some FinTechs dominate and others fail

Connection to This Topic:

- The data flywheel is a platform concept
- Lending marketplaces rely on network effects
- Data advantages compound over time
- Understanding platform economics explains FinTech competition

Before Next Session

- Complete **Notebook NB04** (Credit Scoring Model)
- Think about: what data would *you* want a lender to consider?
- Consider: is algorithmic lending more or less fair than human judgment?

Questions?

Topic 2.3: Data-Driven Finance
Lending, Scoring, and Algorithmic Decision-Making